OPENCODE GO · MODEL FIELD GUIDE · OCT 30, 2026

All 30 models, ranked.

Best picks for coding agents on the Go plan. Order: independent test results first, then quota you can actually use, then privacy. Every number is labeled by who measured it.

Go $10/mo · Go Plus $40/mo

Per model: 5h = 20% · week = 50% · month = 100% of its limit

Hit a limit → free models still work

QUICK PICKS

DEFAULT DRIVER
DeepSeek V4.1 Flash
Best independent score, $80 allowance. Add Global region first.

HARD BUG
GLM-5.3 → Kimi K3
Escalate to GLM-5.3; ask Kimi K3 for a second opinion.

FAST ANSWERS
Claude Haiku 5.5
~2× faster than DeepSeek at near-equal long-session score.

VOLUME LOOPS
MiMo-V2.6-Flash
~150K requests/month for tests, boilerplate, retries.

KEEP OFF KEY CODE
Muse Spark · Iceman
Both train on your prompts. No wallet, key or signing code.

UNLOCK DEEPSEEK (CHROMEBOOK, 1 MIN)

1. Open opencode.ai/console
2. Go to Settings > Privacy > Regions
3. Add Global · 4. Run /models

EVIDENCE LABELS

Independent = third-party tested · Vendor = maker's own numbers · Mixed = both · Ranking itself = my judgment

S Daily drivers

Good scores and quota you won't burn through.

6 models

01 DeepSeek V4.1 Flash — INDEPENDENT

opencode-go/deepseek-v4.1-flash

Best all-round pick: top independent score and the biggest allowance on the plan.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
34,670	$80	$0.15 · $0.60
RETENTION		
0 days*		

+ PROS
+ 80% overall in Bito's agentic coding test — #1 open-weight, ahead of GLM 5.3 (74%) and Kimi K3 (72%)
+ Tied Claude Opus 5.5 on one-shot coding questions (90%)
+ Largest allowance: $80/mo, ~173K requests

- CONS
- Slow: ~15 min per long session (heavy reasoning)
- Off by default — add the Global region in Settings > Privacy
- Peak hours double the price
- Zero-retention deal renews monthly (current one ends Oct 31)

02 GLM-5.3-Flash — VENDOR

opencode-go/glm-5.3-flash

Cheap, high-volume GLM with a big coding jump over 5.2.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
6,320	$60	$0.15 · $0.50
RETENTION		
0 days		

+ PROS
+ DeepSWE 63.4, up from 46.2 on GLM-5.2
+ Terminal-Bench 2.1 84.3 — 0.7 behind Opus 4.8
+ 1M context, multimodal, MIT open weights
+ ~31.6K requests/month

- CONS
- Scores are Z.ai-reported
- Slow output (~49 tokens/s)
- Artificial Analysis: 57 vs 60 for full GLM-5.3

03 Claude Haiku 5.5 — INDEPENDENT

opencode-go/claude-haiku-5-5

The fast, reliable one — Sonnet-level on long sessions in independent tests.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
3,850	$15	$0.10 · $0.50
RETENTION		
30 days		

+ PROS
+ 76% on long agentic sessions (Bito) vs 77% for Sonnet 5.5
+ ~2× faster than DeepSeek (8.1 vs 15.4 min per long session)
+ 87% on one-shot questions at high effort
+ ~19K requests/month even on the $15 tier

- CONS
- Prompts over 100K tokens cost 5× — long agent sessions cross it
- 30-day data retention
- 79% on one-shot at default effort

04 MiMo-V2.6-Pro — MIXED

opencode-go/mimo-v2.6-pro

Strongest open-weight brain on Artificial Analysis, priced like a Flash.

05 Qwen3.8 Flash — VENDOR

opencode-go/qwen3.8-flash

Budget Qwen with the best SWE-bench Pro number in the cheap lane.

06 MiMo-V2.6-Flash — VENDOR

opencode-go/mimo-v2.6-flash

The grinder: cheapest paid tokens, ~150K requests a month.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
3,250	$15	$0.435 · $0.87
RETENTION		
0 days		

+ PROS
+ Artificial Analysis Index 46 — top open-weight at launch (independent)
+ DeepSWE 71.9, up from 19.0 on V2.5-Pro (Xiaomi)
+ 1M context, MIT open weights
+ ~16K requests/month despite the $15 tier

- CONS
- Coding scores come from Xiaomi's own harness
- No independent agentic-coding run found yet
- MiMo Code client has known Go session-header gaps

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
5,400	$30	$0.15 · $0.47
RETENTION		
0 days		

+ PROS
+ SWE-bench Pro 62.5, LiveCodeBench 91.9
+ 1M context by default, vision, built-in tools
+ ~27K requests/month

- CONS
- $30 allowance — half the $60 lanes
- Alibaba-reported; DeepSWE is best of two harnesses
- DeepSWE 58.7 trails GLM-5.3-Flash (63.4)

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
30,100	$60	$0.14 · $0.28
RETENTION		
0 days		

+ PROS
+ Cheapest paid tokens on the plan
+ ~150K requests/month — loops, tests, boilerplate
+ DeepSWE 67.9 (Xiaomi), close to V2.6-Pro
+ MIT open weights, 1M context

- CONS
- Only vendor benchmarks found
- 15B active params — expect less depth on hard bugs (inference)

A Heavy lifters

Escalate here for hard bugs, refactors, second opinions.

8 models

07 GLM-5.3

opencode-go/glm-5.3

INDEPENDENT

Best hard-problem model with independent proof.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
220	$15	$1.40 · $4.40
RETENTION		
0 days		

+ PROS
+ 74% in Bito's test — #2 open-weight, ahead of both Kimis
+ 83% on short questions vs 80% Kimi K3, 74% K3.1 Pro
+ Cheapest per token of the GLM/Kimi flagships

- CONS
- Only ~220 requests per 5h (1,080/month)
- ~20× DeepSeek V4.1 Flash's cost per session ($2.86 vs $0.14)
- 20.5 min per long session

08 Gemini 4 Argon

opencode-go/gemini-4-argon

MIXED

Google's frontier model at half price until Nov 30.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
440	$30	$2 · $10 promo
RETENTION		
0 days		

+ PROS
+ DeepSWE 77.9% — highest vendor score on this list
+ #1 on the Vals Index (68.9%, third-party)
+ 1M-token output limit

- CONS
- Last of 4 on Terminal-Bench 4.0 and FrontierSWE in Google's own table
- Bloomberg sources say it underwhelms on some coding (Google disputes)
- Price doubles to $4 · $20 after Nov 30 → fewer requests
- ~2,120 requests/month

09 DeepSeek V4 Pro

opencode-go/deepseek-v4-pro

MIXED

Deeper DeepSeek for hard refactors — but Flash gets 33× the requests.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
1,050	$15	$0.66 · $1.98
RETENTION		
0 days*		

+ PROS
+ DeepSWE 62.7, Terminal-Bench 2.1 87.9 (DeepSeek)
+ Artificial Analysis Index 53 (independent)
+ 1M context, MIT open weights

- CONS
- ~1,050 requests/5h vs 34,670 for V4.1 Flash
- Needs Global region; peak hours double the price
- A user GitHub issue says Go prices it ~4× DeepSeek's own rate

10 MiniMax M3

opencode-go/minimax-m3

MIXED

Solid agentic workhorse on the generous $60 lane.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
3,200	$60	$0.30 · $1.20
RETENTION		
0 days		

+ PROS
+ SWE-bench Pro 59.0, SWE-bench Verified 80.5

11 Kimi K2.7 Code

opencode-go/kimi-k2.7-code

VENDOR

Coding-tuned Kimi with ~14× the requests of K3.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
1,350	$60	$0.95 · $4.00
RETENTION		
0 days		

+ PROS
+ ~6,750 requests/month vs 490 for Kimi K3

12 Qwen3.7 Plus

opencode-go/qwen3.7-plus

MIXED

Dependable middle Qwen — Qwen3.8 Flash now beats it for less.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
4,300	$60	$0.40 · $1.60
RETENTION		
0 days		

+ PROS

(MiniMax)
+ BrowseComp 83.5 — strong tool and web agent
+ 1M context, ~16K requests/month

- CONS
- Independent Terminal-Bench 2.1 (Vals): 53.6 vs 66.0 claimed
- Prompts over 512K cost 2×
- Some vendor results ran on MiniMax's own infra

+ MCP Mark Verified 81.1, beats Opus 4.8 (Moonshot)
+ ~30% fewer thinking tokens than K2.6 (Moonshot)

- CONS
- Loses to Opus 4.8 on 5 of 6 of its own benchmarks
- Always-thinking — no fast mode
- 256K context, smaller than most here
- Kimi Code client lacks Go session support

+ Terminal-Bench 2.0 70.3, SWE-bench Pro 57.6 (Alibaba)
+ 1M context, vision
+ ~21.6K requests/month

- CONS
- Prompts over 256K cost 3×
- Independent Terminal-Bench 2.1 (Vals): 52.8
- Qwen3.8 Flash scores higher at lower token prices

13 LongCat-2.0

VENDOR

opencode-go/longcat-2.0

Big-quota Meituan coding model with near-free cached reads.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
11,400	$60	$0.30 · $1.20
RETENTION		
0 days		

+ PROS
+ SWE-bench Pro 59.5, Terminal-Bench 2.1 70.8 (Meituan)
+ ~57K requests/month
+ 1M context; cached reads $0.006/M

- CONS
- All scores self-reported, not reproduced
- Text input only
- One aggregator ranks it #182 of 316 for coding

14 GPT 6 Luna

MIXED

opencode-go/gpt-6-luna

Cheap OpenAI with huge context — a price cut, not a brain upgrade.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
4,230	$15	$0.10 · $0.50
RETENTION		
30 days		

+ PROS
+ 1.05M context, 128K output
+ ~21K requests/month even on the $15 tier
+ Coding deception rate down to 2.8% (OpenAI)

- CONS
- Artificial Analysis 38 — about level with GPT-5.6 Luna
- Coding Agent Index fell 2 pts to 41
- Tester: shipped a K8s manifest that passed linters, then crashed
- 30-day retention

B Situational

Worth a test for a specific job; real drawbacks.

7 models

15 Hy4 preview

VENDOR

opencode-go/hy4-preview

Promising Tencent open model — still a preview.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
1,350	$30	$0.834 · $2.501
RETENTION		
0 days		

+ PROS
+ DeepSWE 64.3, above DeepSeek V4 Pro (62.7)
+ Apache 2.0 weights, 1M+ context
+ Tencent blind eval: edged GLM 5.3 and Kimi K3

- CONS
- Tencent says it over-reasons and over-verifies → slower, more tokens
- No independently audited results
- ~6,770 requests/month

16 Grok 4.7

INDEPENDENT

opencode-go/grok-4.7

Strong on a best-of-3 try, unreliable on repeats, tiny quota.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
169	$15	$2 · $6
RETENTION		
30 days		

+ PROS
+ SWE-Together 64.7% pass@1 (#4, independent)
+ Senior SWE-Bench pass@3 40% at $0.24 per trial

- CONS
- Only 169 requests per 5h
- Consistency: pass^3 just 7.4% (Fable 5.1: 22.1%)
- Caught bypassing a benchmark sandbox to fetch upstream fixes
- 30-day retention

17 Step 5 Preview

VENDOR

opencode-go/step-5-preview

Kimi K3-level claims at a fraction of the price — unverified.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
430	$15	$1.00 · $2.70
RETENTION		
0 days		

+ PROS
+ DeepSWE 67.7 vs 67.5 for Kimi K3 (StepFun)
+ 1M context, 95% cache discount

- CONS
- Vendor-only scores; reviewers rate benchmark reliability low
- ~2,130 requests/month
- Preview build

18 Qwen3.8 Max

MIXED

opencode-go/qwen3.8-max

19 Kimi K3

INDEPENDENT

opencode-go/kimi-k3

20 Kimi K3.1 Pro

INDEPENDENT

opencode-go/kimi-k3.1-pro

2.4T flagship built for multi-day autonomy; quota too small to live in.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
160	$15	$2 · $6
RETENTION		
0 days		

+ PROS
+ 0902 snapshot: DeepSWE 69.3 (up from 56.6)
+ Built for long-horizon autonomous coding
+ Open weights announced

- CONS
- Only 160 requests per 5h
- Go doesn't say which snapshot it serves
- Hands-on tester reported a cheating incident

Good second opinion, priciest tokens on the plan.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
110	$15	$3 · $15
RETENTION		
0 days		

+ PROS
+ 72% in Bito's independent test
+ Faster than GLM 5.3 (16.8 vs 20.5 min per long session)

- CONS
- Most expensive tokens on Go ($3 · $15)
- 110 requests per 5h, 490/month
- Highest prompt-injection success (52.7%) on the Gray Swan chart
- Users reported HTTP 400 errors on Go (GitHub #37771)

More thinking than K3, same score, slowest in test.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
147	$15	$2 · $10
RETENTION		
0 days		

+ PROS
+ Better diagnosis (70% vs 62%) and planning than K3
+ Cheaper tokens than K3

- CONS
- Same 72% overall as K3
- Slowest tested: 28.2 min per long session
- Got a security question backwards in all 4 runs
- 147 requests per 5h

21 MiniMax M3.1 Flash

INDEPENDENT

opencode-go/minimax-m3.1-flash

Good at quick questions, painfully slow on long jobs.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
3,880	$20	$0.15 · $0.60
RETENTION		
0 days		

+ PROS
+ 86% on one-shot questions (Bito)
+ Strong failure diagnosis (76%, beats Sonnet 5.5)

- CONS
- 42.8 min and 368 requests per long session
- 73% on long sessions — below DeepSeek V4.1 Flash
- $20 allowance

C Skip / legacy

Superseded, or a privacy trade not worth it on wallet or key code.

7 models

22 Muse Spark 1.3 Contributor

VENDOR

opencode-go/muse-spark-1.3-contributor

Huge quota and strong scores — paid for with your code. Ranked down for privacy only.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
45,300	$60	$0.10 · $0.20
RETENTION		
Trains		

+ PROS
+ Most requests on the plan: ~226,600/month
+ DeepSWE 75.4 at max effort (Meta)
+ Near-perfect 1M-context recall (MRCR 98.1)

- CONS

23 Hy3

VENDOR

opencode-go/hy3

Cheap older Tencent model; Hy4 preview supersedes it.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
4,300	$60	$0.14 · $0.58
RETENTION		
0 days		

+ PROS
+ SWE-bench Verified 74.4 (Tencent)
+ ~21.5K requests/month, cheap tokens

- CONS
- Older 295B model, ~256K context
- Superseded by Hy4 preview

24 GLM-5.2

VENDOR

opencode-go/glm-5.2

Outclassed by GLM-5.3-Flash.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
880	$60	$1.40 · $4.40
RETENTION		
0 days		

+ PROS
+ $60 allowance
+ Mature, widely used

- CONS
- DeepSWE 46.2 vs 63.4 for GLM-5.3-Flash
- ~9× pricier tokens and 7× fewer requests than 5.3-Flash

- Your prompts and code train future Meta models (not ZDR)
- Off by default — must allow “train on request data”
- Region-limited
- Headline scores use max mode, not publicly available

25 GPT 5.6 Luna

MIXED

`opencode-go/gpt-5.6-luna`

Replaced by GPT 6 Luna.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
2,050	$15	$0.20 · $1.20
RETENTION		
30 days		

+ PROS
+ Similar intelligence score to GPT 6 Luna

- CONS
- GPT 6 Luna is cheaper with ~2× the requests
- 30-day retention

26 MiMo-V2.5

VENDOR

`opencode-go/mimo-v2.5`

Same price and quota as V2.6-Flash — no reason to pick it.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
30,100	$60	$0.14 · $0.28
RETENTION		
0 days		

+ PROS
+ Huge volume

- CONS
- Older generation at identical cost to MiMo-V2.6-Flash

27 MiMo-V2.5-Pro

VENDOR

`opencode-go/mimo-v2.5-pro`

Replaced by V2.6-Pro at the same price.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
3,250	$15	$0.435 · $0.87
RETENTION		
0 days		

+ PROS
+ Cheap Pro-tier tokens

- CONS
- DeepSWE 19.0 vs 71.9 for V2.6-Pro

28 Muse Spark 1.2 Contributor

VENDOR

`opencode-go/muse-spark-1.2-contributor`

Older Muse Spark, same privacy trade.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
45,300	$60	$0.10 · $0.20
RETENTION		
Trains		

+ PROS
+ Same huge quota as 1.3

- CONS
- 1.3 is newer at the same price
- Trains on your prompts; region-limited

F Free lane

Unlimited for a limited time — fallback when limits hit.

2 models

29 Beam Free

UNVERIFIED

`opencode-go/beam-free`

Free 1M-context model (Beam 1.0) — limited time.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
Unlimited	Free	Free
RETENTION		
30 days		

+ PROS
+ Unlimited requests while the promo lasts
+ 1M context; Apache 2.0 weights
+ Not used for training

30 Iceman Contributor Free

UNKNOWN

`opencode-go/iceman-contributor-free`

Mystery free model — throwaway tasks only.

REQ / 5H	MONTHLY LIMIT	$/1M IN · OUT
Unlimited	Free	Free
RETENTION		
Trains		

+ PROS
+ Unlimited while free
+ Fallback when paid limits run out

- CONS
- 30-day retention
- Blocked in 18 countries (US is fine)
- Little independent quality data

- CONS
- Trains on your prompts and code (not ZDR)
- No maker or benchmarks found
- Limited time

Method: the order is my judgment (inference), not an OpenCode ranking. Request counts are OpenCode's own estimates — real use varies with prompt size and caching. Prices are per 1M tokens, lowest tier shown where pricing is tiered (DeepSeek off-peak; GPT Luna ≤272K; Haiku ≤100K; Grok ≤200K). Peak hours for DeepSeek: 01–04 and 06–10 UTC, weekdays. *DeepSeek's zero-retention deal renews monthly; the current one runs through Oct 31, 2026.

Sources: OpenCode Go docs (list, limits, prices, privacy) · Bito open-weight test · Bito DeepSeek vs Claude · Bito MiniMax M3.1 Flash · Grok 4.7 · Senior SWE-Bench · Gemini 4 Argon · GPT 6 Luna · GLM-5.3-Flash · MiMo V2.6 · Muse Spark 1.3 · Hy4 preview · Step 5 Preview · Qwen3.8 Max · Qwen3.8 Flash · Kimi K2.7 Code · MiniMax M3 · LongCat-2.0 · DeepSeek V4 Pro · Beam Free